DiD for Big Data in R
Theoretical Background

Bradley Setzler, Penn State

This draft compiled on: October 28, 2022

Notation and Identification (1/4)

Notation and Definition of Treatment:

- i: unit of observation.
  - $\rightarrow$ We will often say "individual".
- t: calendar time. Sample time frame is $\mathcal{T}$.
  - $\rightarrow$ We will often say "year".
- $D_{i,t}$: indicator for currently receiving treatment.
  - $\rightarrow$ Permanent treatment: $D_{i,t} = 1 \implies D_{i,t+1} = 1,\ t \in \mathcal{T}$.
- $G_i \equiv \min\{t : D_{i,t} = 1\}$: time that i first receives treatment.
  - $\rightarrow$ We will often say "cohort" or "onset time".
  - $\rightarrow$ If i never receives treatment, we can write $G_i = \infty$.
  - $\rightarrow$ Note: $D_{i,t} \equiv 1\{t \geq G_i\}$.
- $E_{i,t} \equiv (t - G_i)$: time since first treatment.
  - $\rightarrow$ We will often say "event time".
  - $\rightarrow$ Sometimes we will consider fixing an event time e years after treatment versus b years before treatment.
- Potential outcomes: Let $Y_{i,t}(g)$ denote the outcome that is experienced if treated at g.
- Observed outcome:
  $Y_{i,t} = Y_{i,t}(\infty) + \sum_g 1\{G_i = g\}(Y_{i,t}(g) - Y_{i,t}(\infty))$.

Goal: Identify ATT at an event time. We assume throughout that the goal is to identify the average treatment effect on the treated (ATT) at event time e. Formally, we seek to identify,

$$\text{ATT}_e \equiv \mathbb{E}[Y_{i,t}(g) - Y_{i,t}(\infty)|E_{i,t} = e]$$

The identification challenge is that $\mathbb{E}[Y_{i,t}(\infty)|E_{i,t} = e]$ is a counterfactual object – it is the average outcome that *would have been experienced* by those receiving treatment for e years *if they had not received treatment*.

Decomposition into cohort-specific ATTs. Define,

$$\text{ATT}_{g,e} \equiv \mathbb{E}[Y_{i,g+e}(g) - Y_{i,g+e}(\infty)|G_i = g]$$
$$\omega_{g,e} \equiv \mathbb{E}[G_i = g|E_{i,t} = e]$$

where the cohort shares that have been treated at each event time ($\omega_{g,e}$) are observed. Rearranging terms,

$$\text{ATT}_e = \sum_g \omega_{g,e}\text{ATT}_{g,e}$$

Thus, given e, it is sufficient to identify $\text{ATT}_{g,e}$, $\forall g$ s.t. $\omega_{g,e} > 0$.

Control Group: Define a control group membership indicator $\mathcal{C}_{g,e}(G_i)$. At a minimum, $\mathcal{C}_{g,e}(G_i)=1 \implies G_i > (g+e)$. We may further restrict $\mathcal{C}$ based on context, e.g., some consider the never-treated control group, $\mathcal{C}_{g,e}(G_i)=1\{G_i=\infty\}$.

Assumption 1: Parallel trends.

$$\exists b < 0 \text{ s.t. } \quad \mathbb{E}[Y_{i,g+e}(\infty) - Y_{i,g+b}(\infty)|\mathcal{C}_{g,e}(G_i)=1]$$
$$= \mathbb{E}[Y_{i,g+e}(\infty) - Y_{i,g+b}(\infty)|G_i = g], \ \forall e \geq 0$$

This restricts the relationship between the treated group $G_i = g$ and the control group $\mathcal{C}_{g,e}(G_i)=1$: the change in average outcome for the treated group *would have been the same in the absence of treatment* as that of the control group.

Assumption 2: No anticipation.

$$\exists b < 0 \text{ s.t. } \mathbb{E}[Y_{i,g+b}(g)|G_i = g] = \mathbb{E}[Y_{i,g+b}(\infty)|G_i = g], \ \forall g$$

This restricts *when* the treated cohorts respond to treatment.

Note: Both assumptions need only hold for the event time e and pre-period b chosen by the researcher.

Notation and Identification (4/4)

Difference-in-differences: Define the population estimator,

$\text{DiD}_{g,e} \equiv \mathbb{E}[Y_{i,g+e} - Y_{i,g+b} | G_i = g] - \mathbb{E}[Y_{i,g+e} - Y_{i,g+b} | \mathcal{C}_{g,e}(G_i) = 1]$

It depends only on observed outcomes, not counterfactuals.

Impose parallel trends: By the parallel-trends assumption, we can replace the second expectation as follows:

$\text{DiD}_{g,e} \equiv \mathbb{E}[Y_{i,g+e} - Y_{i,g+b} | G_i = g] - \mathbb{E}[Y_{i,g+e}(\infty) - Y_{i,g+b}(\infty) | G_i = g]$

Impose no anticipation: By the no-anticipation assumption, we can cancel out the two terms involving $Y_{i,g+b}$:

$\text{DiD}_{g,e} \equiv \mathbb{E}[Y_{i,g+e} | G_i = g] - \mathbb{E}[Y_{i,g+e}(\infty) | G_i = g] = \text{ATT}_{g,e}$

Thus, we have proven that $\text{DiD}_{g,e} = \text{ATT}_{g,e}$ if the parallel-trends and no-anticipation assumptions hold for the pair (g, e).

Result: If parallel-trends and no-anticipation hold $\forall g$ s.t. $\omega_{g,e} > 0$,

$$\text{ATT}_e = \sum_{g:\, w_{g,e} > 0} \omega_{g,e} \text{DiD}_{g,e}$$

Estimator based on averages. Replacing population means with sample means, the package implements the following DiD:

$$\mathrm{DiD}_{g,e} = \underbrace{\mathbb{E}[Y_{i,g+e} - Y_{i,g+b}|G_i = g]}_{\textbf{Difference for treated group}} - \underbrace{(\mathbb{E}[Y_{i,g+e} - Y_{i,g+b}|\mathcal{C}_{g,e}(G_i) = 1])}_{\textbf{Difference for control group}}$$

where, following Callaway and Sant'Anna (2021), we require that parallel-trends and no-anticipation holds for one of these 3 possible control groups:

$$\text{``all''} \quad \mathcal{C}_{g,e}(G_i) = 1\{G_i > (g+e)\}$$
$$\text{``future-treated''} \quad \mathcal{C}_{g,e}(G_i) = 1\{G_i > (g+e) \ \& \ G_i < \infty\}$$
$$\text{``never-treated''} \quad \mathcal{C}_{g,e}(G_i) = 1\{G_i = \infty\}$$

Researcher choices. The DiD researcher must make 3 choices:

1. What is the range of event times e for which you would like ATT_e estimates? Default: $e = -5, ..., 5$.
2. Which pre-period should be the base? Default: $b = -1$.
3. Which of the 3 control selections $\mathcal{C}$ to use? Default: "all".

Estimator used by DiD for Big Data (2/4)

Fix some (e, b, g). Consider testing $\text{ATT}_{g,e} = 0$.

Notation: Consider treatment group $\mathcal{T} \equiv 1\{G_i = g\}$ and control group $\mathcal{C}$, which could be any of the three options above.

Define within-i differences $A_i \equiv Y_{i,g+e} - Y_{i,g+b}, i \in \mathcal{T}$ with mean $\mu_A \equiv \mathbb{E}[A_i | i \in \mathcal{T}]$, and $B_i \equiv Y_{i,g+e} - Y_{i,g+b}, i \in \mathcal{C}$ with mean $\mu_B \equiv \mathbb{E}[B_i | i \in \mathcal{C}]$, where subscripts are dropped if unambiguous.

Hypothesis Testing: Since $\text{ATT}_{g,e} \equiv \mu_A - \mu_B$, consider,

$$\text{Test statistic: } \text{DiD}_{g,e} = \overline{A} - \overline{B}, \quad \text{Null } H_0 : \mu_A - \mu_B = 0$$

Central Limit Theorem: Denote the population variances by $\sigma_A^2 \equiv \text{Var}[A_i | i \in \mathcal{T}]$ and $\sigma_B^2 \equiv \text{Var}[B_i | i \in \mathcal{C}]$. By the CLT under the null, and that the samples are drawn independently across i,

$$\overline{A} \sim_d \mathcal{N}\left(\mu_A, \ \sigma_A^2 / N_A\right), \quad \overline{B} \sim_d \mathcal{N}\left(\mu_B, \ \sigma_B^2 / N_B\right)$$

$$\implies \text{DiD}_{g,e} = \left(\overline{A} - \overline{B}\right) \sim_d \mathcal{N}\left(0, \ \sigma_A^2 / N_A + \sigma_B^2 / N_B\right)$$

Thus, $\text{SE}(\text{DiD}_{g,e}) = \sqrt{\sigma_A^2 / N_A + \sigma_B^2 / N_B}$. The empirical counterpart is trivial to compute (e.g. no matrix inversion needed).

Estimator used by DiD for Big Data (3/4)

We usually are not interested in the ATT for a specific cohort g. Instead, we are usually interested in the ATT that is e years after treatment, averaging across cohorts.

Average Effects by Event Time. Let ω_g denote the share of treated units in cohort g. Since treated cohorts are independent of one another at a given e, CLT implies,

$$\text{DiD}_e = \sum_{\mathcal{G}} \omega_g \text{DiD}_{g,e}, \quad \text{SE}(\text{DiD}_e) = \sqrt{\sum_{g \in \mathcal{G}} \omega_g^2 \text{Var}(\text{DiD}_{g,e})}$$

Average across Event Times: Similarly, we are often interested in an average of DiD_e across event times. Letting $\mathcal{E}$ denote a set of event times (e.g. $\mathcal{E} = \{1, 2, 3\}$), we are typically interested in the equally-weighted average, which has the following SE:

$$\text{DiD}_{\mathcal{E}} = \frac{1}{|\mathcal{E}|} \sum_{\mathcal{E}} \text{DiD}_e, \quad \text{SE}(\text{DiD}_{\mathcal{E}}) = \frac{1}{|\mathcal{E}|} \sqrt{\sum_{e \in \mathcal{E}} \text{Var}(\text{DiD}_e)}$$

The R package provides all possible DiD_e with their SEs by default, and user-friendly options to estimate various $\text{DiD}_{\mathcal{E}}$ and their SEs.

Estimator used by DiD for Big Data (4/4)

Event study parameters: For $\mathcal{H} \in \{\mathcal{T}, \mathcal{C}\}$, one may wish to plot averages across event times:

$$\bar{Y}_{g,e}^{\mathcal{H}} = \sum_{i \in \mathcal{H}_g} Y_{i,g+e}, \; \mathsf{SE}(\bar{Y}_{g,e}^{\mathcal{H}}) = \sqrt{\frac{\sigma_{\mathcal{H},g,e}^2}{|\mathcal{H}_g|}}, \; \sigma_{\mathcal{H},g,e}^2 = \mathsf{Var}[Y_{i,g+e} | i \in \mathcal{H}_g]$$

$$\bar{Y}_e^{\mathcal{H}} = \sum_{g \in \mathcal{G}} \omega_g \bar{Y}_{g,e}^{\mathcal{H}}, \; \mathsf{SE}(\bar{Y}_e^{\mathcal{H}}) = \sqrt{\sum_g \omega_g^2 \mathsf{Var}(\bar{Y}_{g,e}^{\mathcal{H}})},$$

The package provides these means and SEs by default.

Plots: The package also provides automated plots, both for presenting the event study parameters and for presenting the ATT estimates. These can be plotted by (g, e) or by e (average over g).